

## **Yuxuan (Wayne) Wang**

[https://waynewangyuxuan.github.io/personal\\_site/](https://waynewangyuxuan.github.io/personal_site/) || waw009@ucsd.edu

### **EDUCATION SUMMARY:**

**University of California, San Diego**, Jacobs School of Engineering, San Diego, CA  
Master of Science, Computer Science September 2025 - June 2027

**New York University**, Tandon School of Engineering, Brooklyn, NY  
Bachelor of Science, Computer Science & Math September 2021 - May 2025

### **TECHNICAL SKILLS:**

**C/C++ || Java || Python || Linux** (Kernel, Shell/Bash) || **Multi-Threaded/Process Programming** (Threads, Memory Management, IPC) || **Parallel Computing** (CUDA) || **Agent/Workflow Development** (LangGraph, n8n, Dify) || **Information Retrieval** (query processing, dense/sparse re-ranking) || **Computer Networking** (TCP/IP, OSI Model, Sockets) ||

### **EXPERIENCES:**

**Data Engineering Intern - ByteDance** June 2025 – September 2025

Global Monetization Product and Technology - TikTok Ads Diagnosis  
(Java, Python, SQL, Kafka, Hive, Spark, Big Data, TikTok Ads)

- Migrated 20+ ads analytics pipelines (Kafka→Hive/HDFS; ClickHouse→Doris), refactoring schemas/HiveSQL and shipping a Java engine adapter; built automated QA (schema/volume diffs, query-parity, p50/p95 latency tracking) to gate multi-DC cutovers.
- Designed and shipped Intent-vs-Actual Budget Delta metric; re-architected 2 primary and upgraded 3 downstream pipelines (~150 TB backfill) to align budget caliber across delivery stages; exposed results via materialized views/APIs for AdDebug.
- Hardened reliability & performance with idempotent backfills, partition-safe replays, and chat-ops alerts; reduced rollout incidents and query drift, and authored runbooks to onboard partner teams.

**Research - New York University** June 2024 – May 2025

**Evaluation of Graph-Based Vocabulary Mismatch Solution in Information Retrieval**

Supervised by Professor Torsten Suel @ New York University, Tandon School of Engineering  
(Information Retrieval, Search Engine, Query Processing, Database, Linux, HPC, Python)

- Conducted a systematic study on impact of seed quality graph-based expansion in LADR (Lexically-Accelerated Dense Retrieval); designed a plug-and-play retrieval pipeline supporting multiple sparse retrieval models as initial seed sets.
- Implemented a graph expansion module for candidate documents using KNN/HNSW; incorporated a vector dimension masking mechanism (PRFDIME) that constructs semantic masks via pseudo-relevance feedback to enhance query representation and improve reranking effectiveness.
- Demonstrated that reranking using only the top-3000 graph-expanded passages via Bi-Encoder/Cross-Encoder models can achieve comparable performance to full-corpus reranking, while improving recall and reducing computational cost.

**Data Engineering Intern - CITIC Poly Fund (Guangzhou)** June 2023 - August 2023

(NLP-based News ETL, Web Scraping, Python, AWS, NoSQL)

- Built an end-to-end investment-intelligence platform for the chip/EDA sector—ingesting Wind + multi-site crawlers into a Kafka→Lambda/DynamoDB/S3 pipeline with de-dup, structured field extraction (amount/round/investors), and company-profile enrichment—and delivered an insights layer (event/keyword mining, BERT sentiment, T5 summarization, weekly trends) with a Streamlit dashboard that replaced manual spreadsheets with a department-wide weekly brief.

### **PROJECTS:**

**TripPlanner** March 2024 - May 2025

- Built an AWS serverless travel assistant—owning route optimization/mapping and real-time notifications: integrated Amazon Location Service + OSRM for multi-stop routing (TSP heuristic, K-d tree), produced GeoJSON with static map renders to S3, added Redis caching to de-dupe/accelerate repeat routes, and delivered user-tunable alerts via Step Functions/EventBridge/SNS with OpenWeatherMap integration; exposed via API Gateway/Lambda with DynamoDB persistence and an Amplify front end.

**ShadowDash** September 2024 - Jan 2025

- Built a C++ DSL that bypasses Ninja's .ninja text parsing by constructing the build DAG in-memory (rule/build/pool); validated on zlib and LLVM, cutting parse time by ~40% (zlib, -O3) and ~2× (LLVM, -O1), with macros & user-defined literals for concise manifests and direct integration with Ninja state.

